

▼ Deep Learning: Homework 1

Author: Rudy Martinez

Date: 8/28/2021

Google Colab Link: https://colab.research.google.com/drive/18foX8skV_cUcMrFoC9OKSEHc-5xmXwmf?usp=sharing

▼ Libraries

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sb
```

▼ Part 1

(1) Import the iris flowers dataset using `pandas.read_csv()` with the following URL link (10pt); The DataFrame must have the following column names: `sepal length (cm)`, `sepal width (cm)`, `petal length (cm)`, `petal width (cm)`, and `class`; (5pt) Print the first 5 rows of the DataFrame (5pt).

```
data = pd.read_csv("http://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data",
                  names=['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)', 'class'])

iris_df = pd.DataFrame(data)

iris_df.head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

(2) Summarize the dataset

- Print out a concise summary of the DataFrame using `.info()` and the shape of the DataFrame (5 pt)

```
print(iris_df.info())
print(iris_df.shape)
```

```
print("\n\n Shape:", iris_df.shape)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sepal length (cm)      150 non-null   float64
1   sepal width (cm)       150 non-null   float64
2   petal length (cm)      150 non-null   float64
3   petal width (cm)       150 non-null   float64
4   class                  150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
None

Shape: (150, 5)
```

- Print out the statistics of the continuous columns using .describe() (i.e., the four attribute columns) (5 pt)

```
print(iris_df[0:3].describe())
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
count	3.0	3.000000	3.000000	3.000000e+00
mean	4.9	3.233333	1.366667	2.000000e-01
std	0.2	0.251661	0.057735	3.399350e-17
min	4.7	3.000000	1.300000	2.000000e-01
25%	4.8	3.100000	1.350000	2.000000e-01
50%	4.9	3.200000	1.400000	2.000000e-01
75%	5.0	3.350000	1.400000	2.000000e-01
max	5.1	3.500000	1.400000	2.000000e-01

- Print out the number of rows that belong to each class (5 pt)

```
iris_pivot = pd.pivot_table(iris_df, values = "sepal length (cm)", index = "class", aggfunc = len, margins = True)
iris_pivot.columns = ['# of Rows']
iris_pivot
```

# of Rows	
class	
Iris-setosa	50.0
Iris-versicolor	50.0
Iris-virginica	50.0
All	150.0

(3) Data Visualization

- Separate out the first four columns of the original DataFrame into a new DataFrame and print out the first 5 rows of the new DataFrame (5

```
iris_subset = iris_df[['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']]
iris_subset.head()
```

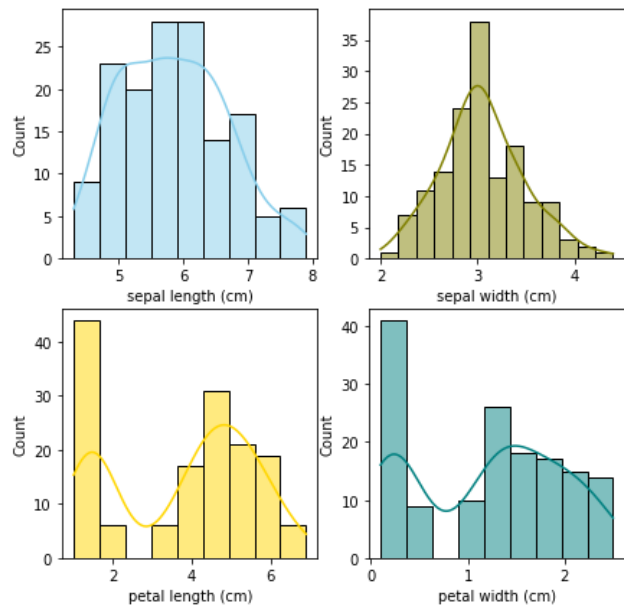
	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

- Univariate Plots: plot a histogram for each column of the new DataFrame (5 pt)

```
fig, axs = plt.subplots(2, 2, figsize=(7, 7))
```

```
sb.histplot(iris_subset, x="sepal length (cm)", kde=True, color="skyblue", ax=axs[0, 0])
sb.histplot(iris_subset, x="sepal width (cm)", kde=True, color="olive", ax=axs[0, 1])
sb.histplot(iris_subset, x="petal length (cm)", kde=True, color="gold", ax=axs[1, 0])
sb.histplot(iris_subset, x="petal width (cm)", kde=True, color="teal", ax=axs[1, 1])
```

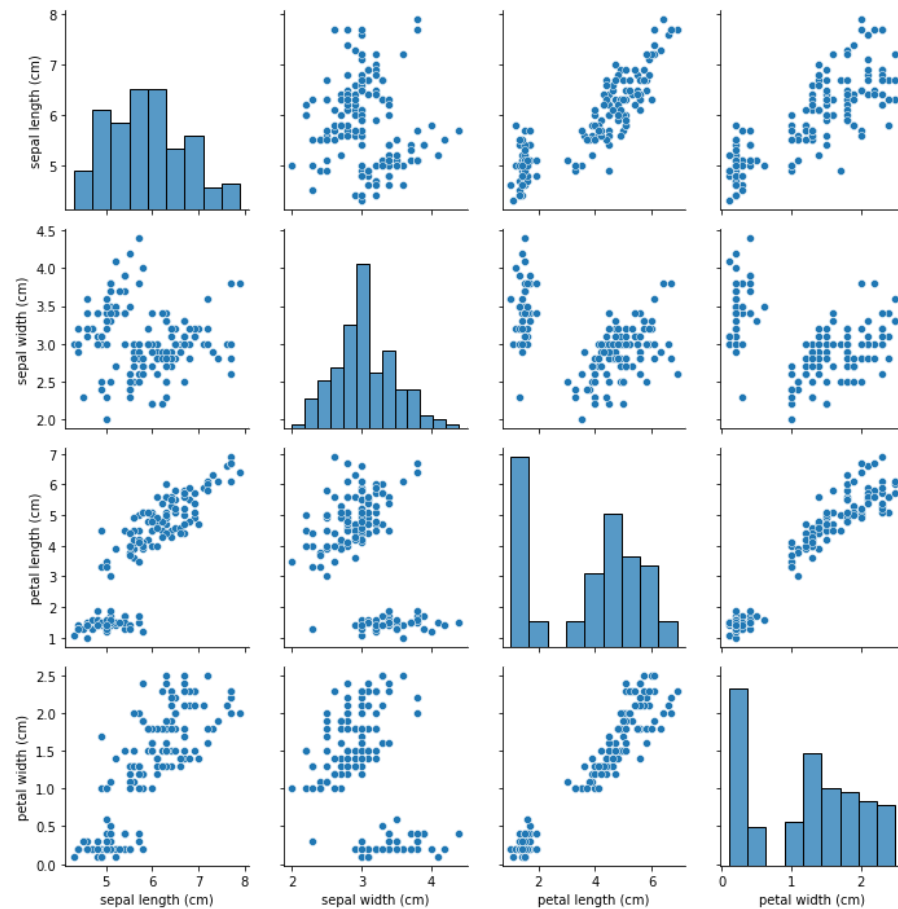
```
plt.show()
```



- Multivariate Plots: plot a scatter plot for each pair of the columns of the new DataFrame using the `pandas.plotting.scatter_matrix` function (5 pt)

```
multi_plot = sb.PairGrid(iris_subset)
multi_plot.map_diag(sb.histplot)
multi_plot.map_offdiag(sb.scatterplot)
```

<seaborn.axisgrid.PairGrid at 0x7f23b4ee9890>



▼ Part 2

(1) Import the `Census Income (Adult)` dataset using Pandas, use the 14 attribute names (i.e., `age`, `workclass`,, `native-country`) as explained in the dataset description as the first 14 column names and `salary` as the last column name (5 pt), view the strings `'?'`, `' ?'`, `' ? '`, or `' ? ' ? '`

'?' as the missing values and replace them with NaN (the default missing value marker in Pandas) (10 pt), and print out the first five rows of the DataFrame. (5 pt)

```
census_data = pd.read_csv("http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data",
                           names = ['age', 'workclass', 'fnlwgt', 'education', 'education-num', 'marital-status',
                                     'occupation', 'relationship', 'race', 'sex', 'capital-gain', 'capital-loss',
                                     'hours-per-week', 'native-country', 'salary'],
                           index_col = False)
```

```
missing_values = ['?', ' ? ', '? ', ' ? ']
```

```
census_df = census_data.replace(missing_values, np.nan)
census_df.head()
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Married-civ-spouse
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Married-civ-spouse
4	28	Private	339400	Bachelors	13	Married-civ-spouse	Prof-specialty	Married-civ-spouse

(2) Dataset checking and cleaning

- Print out a concise summary of the DataFrame and observe if null values exist in each column of the DataFrame by checking the summary (10 pt)

```
census_df.isnull().sum()
```

```
age          0
workclass    1836
fnlwgt       0
education    0
education-num 0
marital-status 0
occupation   1843
relationship 0
race         0
sex          0
capital-gain 0
capital-loss 0
hours-per-week 0
native-country 583
```

```

salary          0
dtype: int64

```

- Filter out the rows that contain missing values and print them out (10 pt)

```

null_rows = census_df[census_df.isna().any(axis=1)]
print(null_rows)

```

```

   age  workclass  fnlwgt  ...  hours-per-week  native-country  salary
14   40   Private  121772  ...             40             NaN    >50K
27   54         NaN  180211  ...             60             South  >50K
38   31   Private   84154  ...             38             NaN    >50K
51   18   Private  226956  ...             30             NaN    <=50K
61   32         NaN  293936  ...             40             NaN    <=50K
...   ...       ...      ...      ...      ...      ...      ...
32530  35         NaN  320084  ...             55  United-States    >50K
32531  30         NaN   33811  ...             99  United-States    <=50K
32539  71         NaN  287372  ...             10  United-States    >50K
32541  41         NaN  202822  ...             32  United-States    <=50K
32542  72         NaN  129912  ...             25  United-States    <=50K

```

```

[2399 rows x 15 columns]

```

- Drop the rows of the DataFrame with missing values using `.dropna()` and observe if null values still exist in each column by checking the concise summary again (10 pt)

```

census_new = census_df.dropna(axis=1)
census_new.isnull().sum()

```

```

age          0
fnlwgt       0
education    0
education-num 0
marital-status 0
relationship 0
race         0
sex          0
capital-gain 0
capital-loss 0
hours-per-week 0
salary       0
dtype: int64

```

