

IS 6733

Deep Learning on Cloud Platforms

Lecture 2b

Python Tutorial - Pandas

Dr. Yuanxiong Guo
Department of Information Systems and Cyber Security



Python Checklist in Machine Learning

- Essential libraries and tools in data science
 - Jupyter Notebook/Colab
 - NumPy
 - Pandas
 - Matplotlib
 - Scikit-Learn
 - Keras/TensorFlow

Pandas

- *[pandas] is derived from the term "panel data", an econometrics term for data sets that include observations over multiple time periods for the same individuals. — Wikipedia*
- This tool is essentially your data's home. Through pandas, you get acquainted with your data by cleaning, transforming, and analyzing it.

What's Pandas for?

- Say you want to explore a dataset stored in a CSV on your computer. Pandas will extract the data from that CSV into a DataFrame — a table, basically — then let you do things like:
 - Calculate statistics and answer questions about the data, like
 - What's the average, median, max, or min of each column?
 - Does column A correlate with column B?
 - What does the distribution of data in column C look like?
 - Clean the data by doing things like removing missing values and filtering rows or columns by some criteria
 - Visualize the data with help from Matplotlib. Plot bars, lines, histograms, bubbles, and more.
 - Store the cleaned, transformed data back into a CSV, other file or database

How dose Pandas fit into the data science toolkit?

- Pandas is built on top of the **NumPy** package, meaning a lot of the structure of NumPy is used or replicated in Pandas. Data in pandas is often used to feed statistical analysis in **SciPy**, plotting functions from **Matplotlib**, and machine learning algorithms in **Scikit-learn**.
- Jupyter Notebooks offer a good environment for using pandas to do data exploration and modeling, but pandas can also be used in text editors just as easily.

When should you start using pandas?

- If you do not have any experience coding in Python, then you should stay away from learning pandas until you do.
- You don't have to be at the level of the software engineer, but you should be adept at the basics, such as lists, tuples, dictionaries, functions, and iterations.
- Familiarize yourself with **NumPy** due to the similarities mentioned above.

Install and Import

- `conda install pandas`
- `pip install pandas`
- `!pip install pandas` (in a Jupyter notebook)
 - The `!` at the beginning runs cells as if they were in a terminal.
 - Common libraries such as `pandas`, `numpy`, `matplotlib` have already been installed in Colab. Therefore, you can directly import and use them.
- `import pandas as pd`

Core Components of Pandas

- Two primary two components of Pandas: **Series** and **DataFrame**
- A **Series** is essentially a column, and a **DataFrame** is a multi-dimensional table made up of a collection of Series.
- A Series is an analog of a one-dimensional array with flexible indices.
- A DataFrame is an analog of a two-dimensional array with both flexible row indices and flexible column names.

Series			Series			DataFrame		
	apples			oranges			apples	oranges
0	3	+	0	0	=	0	3	0
1	2		1	3		1	2	3
2	0		2	7		2	0	7
3	1		3	2		3	1	2

Creating Series from Scratch

- A Pandas Series is a one-dimensional array of indexed data. It can be created from a list or array as follows:

```
In [2]: data = pd.Series([0.25, 0.5, 0.75, 1.0])  
data
```

```
Out[2]: 0    0.25  
        1    0.50  
        2    0.75  
        3    1.00  
        dtype: float64
```

- As we see in the output, the *Series* wraps both a sequence of values and a sequence of indices, which we can access with the `values` and `index` attributes.

```
In [3]: data.values
```

```
Out[3]: array([ 0.25,  0.5 ,  0.75,  1.  ])
```

```
In [4]: data.index
```

The index is an array-like object of type `pd.Index`.

```
Out[4]: RangeIndex(start=0, stop=4, step=1)
```

```
In [5]: data[1]
```

```
Out[5]: 0.5
```

```
In [6]: data[1:3]
```

Like with a NumPy array, data can be accessed by the associated index via the familiar Python square-bracket notation

```
Out[6]: 1    0.50  
        2    0.75  
        dtype: float64
```

Series as generalized NumPy array

- While the NumPy Array has an *implicitly defined* integer index used to access the values, the Pandas Series has an *explicitly defined* index associated with the values.

```
In [7]: data = pd.Series([0.25, 0.5, 0.75, 1.0],  
                        index=['a', 'b', 'c', 'd'])  
data
```

we can use strings as an index

```
Out[7]: a    0.25  
       b    0.50  
       c    0.75  
       d    1.00  
       dtype: float64
```

```
In [8]: data['b']
```

```
Out[8]: 0.5
```

```
In [9]: data = pd.Series([0.25, 0.5, 0.75, 1.0],  
                        index=[2, 5, 3, 7])  
data
```

We can use non-contiguous or non-sequential indices.

```
Out[9]: 2    0.25  
       5    0.50  
       3    0.75  
       7    1.00  
       dtype: float64
```

```
In [10]: data[5]
```

```
Out[10]: 0.5
```

Series as specialized dictionary

- A dictionary is a structure that maps arbitrary keys to a set of arbitrary values, and a *Series* is a structure which maps typed keys to a set of typed values.

```
In [11]: population_dict = {'California': 38332521,
                             'Texas': 26448193,
                             'New York': 19651127,
                             'Florida': 19552860,
                             'Illinois': 12882135}
population = pd.Series(population_dict)
population
```

We can also construct a Series object directly from a Python dictionary.

```
Out[11]: California    38332521
         Florida      19552860
         Illinois     12882135
         New York     19651127
         Texas        26448193
         dtype: int64
```

```
In [12]: population['California']
```

Typical dictionary-style item access can be performed.

```
Out[12]: 38332521
```

```
In [13]: population['California':'Illinois']
```

Unlike a dictionary, though, the Series also supports array-style operations such as slicing.

```
Out[13]: California    38332521
         Florida      19552860
         Illinois     12882135
         dtype: int64
```

Constructing *Series* objects

```
>>> pd.Series(data, index=index)
```

where `index` is an optional argument, and `data` can be one of many entities.

For example, `data` can be a list or NumPy array, in which case `index` defaults to an integer sequence:

```
In [14]: pd.Series([2, 4, 6])
```

```
Out[14]: 0    2  
         1    4  
         2    6  
         dtype: int64
```

`data` can be a dictionary

```
In [16]: pd.Series({'a': 2, 'b': 1, 'c': 3})
```

```
Out[16]: a    2  
         b    1  
         c    3  
         dtype: object
```

`data` can be a scalar, which is repeated to fill the specified index:

```
In [15]: pd.Series(5, index=[100, 200, 300])
```

```
Out[15]: 100    5  
         200    5  
         300    5  
         dtype: int64
```

In each case, the index can be explicitly set if a different result is preferred:

```
In [17]: pd.Series({'a': 2, 'b': 1, 'c': 3}, index=[3, 2])
```

```
Out[17]: 3    c  
         2    a  
         dtype: object
```

Notice that in this case, the `Series` is populated only with the explicitly identified keys.

DataFrame as generalized NumPy array

- If a Series is an analog of a one-dimensional array with flexible indices, a DataFrame is an analog of a two-dimensional array **with both flexible row indices and flexible column names.**

```
In [18]: area_dict = {'California': 423967, 'Texas': 695662, 'New York': 141297,
                    'Florida': 170312, 'Illinois': 149995}
area = pd.Series(area_dict)
area
```

```
Out[18]: California    423967
Florida      170312
Illinois     149995
New York     141297
Texas        695662
dtype: int64
```

```
In [19]: states = pd.DataFrame({'population': population,
                              'area': area})
states
```

```
Out[19]:
```

	area	population
California	423967	38332521
Florida	170312	19552860
Illinois	149995	12882135
New York	141297	19651127
Texas	695662	26448193

```
In [20]: states.index
```

```
Out[20]: Index(['California', 'Florida', 'Illinois', 'New York', 'Texas'], dtype=object)
```

```
In [21]: states.columns
```

```
Out[21]: Index(['area', 'population'], dtype='object')
```

DataFrame as specialized dictionary

- Where a dictionary maps a key to a value, a *DataFrame* maps a column name to a *Series* of column data.

```
In [22]: states['area']
```

```
Out[22]: California    423967  
Florida      170312  
Illinois     149995  
New York     141297  
Texas        695662  
Name: area, dtype: int64
```

- Potential point of confusion:** in a two-dimensional *NumPy* array, `data[0]` will return the first row. For a *DataFrame*, `data['col0']` will return the first column.

Constructing *DataFrame* objects

```
In [23]: pd.DataFrame(population, columns=['population'])
```

```
Out[23]:
```

	population
California	38332521
Florida	19552860
Illinois	12882135
New York	19651127
Texas	26448193

From a single Series object: a DataFrame is a collection of Series objects, and a single-column DataFrame can be constructed from a single Series.

```
In [24]: data = [{ 'a': i, 'b': 2 * i }  
              for i in range(3)]  
pd.DataFrame(data)
```

```
Out[24]:
```

	a	b
0	0	0
1	1	2
2	2	4

From a list of dicts: any list of dictionaries can be made into a DataFrame.

```
In [26]: pd.DataFrame({ 'population': population,  
                       'area': area })
```

```
Out[26]:
```

	area	population
California	423967	38332521
Florida	170312	19552860
Illinois	149995	12882135
New York	141297	19651127
Texas	695662	26448193

From a dictionary of *Series* objects: a DataFrame can be constructed from a dictionary of Series objects

```
In [27]: pd.DataFrame(np.random.rand(3, 2),  
                      columns=['foo', 'bar'],  
                      index=['a', 'b', 'c'])
```

```
Out[27]:
```

	foo	bar
a	0.865257	0.213169
b	0.442759	0.108267
c	0.047110	0.905718

From a two-dimensional NumPy array: we can create a DataFrame with any specified column and index names. If omitted, an integer index will be used for each.

Data Indexing and Selection

- How to assess and modify values in Pandas *Series* and *DataFrame* objects?
- Similar to NumPy in most cases, but there are a few quirks to be aware of.
- Remember that both *Series* and *DataFrame* objects can be thought as generalized NumPy array or specialized dictionary. Therefore, their operations are similar in many ways.

Data Selection in Series

- A *Series* object acts in many ways like a one-dimensional NumPy array, and in many ways like a standard Python dictionary.
- Series as dictionary

```
In [1]: import pandas as pd
data = pd.Series([0.25, 0.5, 0.75, 1.0],
                 index=['a', 'b', 'c', 'd'])
data

Out[1]: a    0.25
        b    0.50
        c    0.75
        d    1.00
        dtype: float64
```

```
In [2]: data['b']

Out[2]: 0.5
```

We can also use dictionary-like Python expressions and methods to examine the keys/indices and values:

```
In [3]: 'a' in data
```

```
Out[3]: True
```

```
In [4]: data.keys()
```

```
Out[4]: Index(['a', 'b', 'c', 'd'], dtype='object')
```

```
In [5]: list(data.items())
```

```
Out[5]: [('a', 0.25), ('b', 0.5), ('c', 0.75), ('d', 1.0)]
```

You can extend a Series by assigning to a new index value:

```
In [6]: data['e'] = 1.25
data
```

```
Out[6]: a    0.25
        b    0.50
        c    0.75
        d    1.00
        e    1.25
        dtype: float64
```

Data Selection in Series

- Series as one-dimensional array

```
In [7]: # slicing by explicit index  
data['a':'c']
```

When slicing with an explicit index (i.e., `data['a':'c']`), the final index is *included* in the slice.

```
Out[7]: a    0.25  
       b    0.50  
       c    0.75  
       dtype: float64
```

```
In [8]: # slicing by implicit integer index  
data[0:2]
```

when slicing with an implicit index (i.e., `data[0:2]`), the final index is *excluded* from the slice.

```
Out[8]: a    0.25  
       b    0.50  
       dtype: float64
```

```
In [9]: # masking  
data[(data > 0.3) & (data < 0.8)]
```

```
Out[9]: b    0.50  
       c    0.75  
       dtype: float64
```

```
In [10]: # fancy indexing  
data[['a', 'e']]
```

```
Out[10]: a    0.25  
        e    1.25  
        dtype: float64
```

To avoid the potential confusion in explicit and implicit indices, Pandas provides some special *indexer* attributes that explicitly expose certain indexing schemes.

Data Selection in Series

- To avoid the potential confusion in explicit and implicit indices, Pandas provides some special *indexer* attributes (*loc*, *iloc*, and *ix*) that explicitly expose certain indexing schemes.
- The *loc* attribute allows indexing and slicing that always references the explicit index, while the *iloc* attribute allows indexing and slicing that always references the implicit Python-style index. A third indexing attribute, *ix*, is a hybrid of the two, and for Series objects is equivalent to standard *[]*-based indexing.

```
In [11]: data = pd.Series(['a', 'b', 'c'], index=[1, 3, 5])  
data
```

```
Out[11]: 1    a  
         3    b  
         5    c  
dtype: object
```

```
In [12]: # explicit index when indexing  
data[1]
```

```
Out[12]: 'a'
```

```
In [13]: # implicit index when slicing  
data[1:3]
```

```
Out[13]: 3    b  
         5    c  
dtype: object
```

```
In [14]: data.loc[1]
```

```
Out[14]: 'a'
```

```
In [15]: data.loc[1:3]
```

```
Out[15]: 1    a  
         3    b  
dtype: object
```

```
In [16]: data.iloc[1]
```

```
Out[16]: 'b'
```

```
In [17]: data.iloc[1:3]
```

```
Out[17]: 3    b  
         5    c  
dtype: object
```

Data Selection in DataFrame

- A *DataFrame* acts in many ways like a two-dimensional or structured array, and in other ways like a dictionary of *Series* structures sharing the same index.
- *DataFrame* as a dictionary

```
In [18]: area = pd.Series({'California': 423967, 'Texas': 695662,
                          'New York': 141297, 'Florida': 170312,
                          'Illinois': 149995})
pop = pd.Series({'California': 38332521, 'Texas': 26448193,
                'New York': 19651127, 'Florida': 19552860,
                'Illinois': 12882135})
data = pd.DataFrame({'area':area, 'pop':pop})
data
```

```
Out[18]:
```

	area	pop
California	423967	38332521
Florida	170312	19552860
Illinois	149995	12882135
New York	141297	19651127
Texas	695662	26448193

Adding a new column:

```
In [23]: data['density'] = data['pop'] / data['area']
data
```

```
Out[23]:
```

	area	pop	density
California	423967	38332521	90.413926
Florida	170312	19552860	114.806121
Illinois	149995	12882135	85.883763
New York	141297	19651127	139.076746
Texas	695662	26448193	38.018740

The columns of the *DataFrame* can be accessed via dictionary-style indexing of the column name:

```
In [19]: data['area']
```

```
Out[19]: California    423967
Florida      170312
Illinois     149995
New York     141297
Texas        695662
Name: area, dtype: int64
```

We can use attribute-style access **only** with column names that are strings:

```
In [20]: data.area
```

```
Out[20]: California    423967
Florida      170312
Illinois     149995
New York     141297
Texas        695662
Name: area, dtype: int64
```

Data Selection in DataFrame

- DataFrame as two-dimensional array

Use the values attribute to examine the raw underlying data array

```
In [24]: data.values
```

```
Out[24]: array([[ 4.23967000e+05,  3.83325210e+07,  9.04139261e+01],
 [ 1.70312000e+05,  1.95528600e+07,  1.14806121e+02],
 [ 1.49995000e+05,  1.28821350e+07,  8.58837628e+01],
 [ 1.41297000e+05,  1.96511270e+07,  1.39076746e+02],
 [ 6.95662000e+05,  2.64481930e+07,  3.80187404e+01]])
```

Can transpose the full DataFrame to swap rows and columns

```
In [25]: data.T
```

```
Out[25]:
```

	California	Florida	Illinois	New York	Texas
area	4.239670e+05	1.703120e+05	1.499950e+05	1.412970e+05	6.956620e+05
pop	3.833252e+07	1.955286e+07	1.288214e+07	1.965113e+07	2.644819e+07
density	9.041393e+01	1.148061e+02	8.588376e+01	1.390767e+02	3.801874e+01

```
In [28]: data.iloc[:3, :2]
```

```
Out[28]:
```

	area	pop
California	423967	38332521
Florida	170312	19552860
Illinois	149995	12882135

```
In [29]: data.loc['Illinois', : 'pop']
```

```
Out[29]:
```

	area	pop
California	423967	38332521
Florida	170312	19552860
Illinois	149995	12882135

passing a single index to an array accesses a row:

```
In [26]: data.values[0]
```

```
Out[26]: array([ 4.23967000e+05,  3.83325210e+07,  9.04139261e+01])
```

passing a single "index" to a DataFrame accesses a column

```
In [27]: data['area']
```

```
Out[27]: California    423967
Florida      170312
Illinois     149995
New York     141297
Texas        695662
Name: area, dtype: int64
```

```
In [30]: data.ix[:3, : 'pop']
```

```
Out[30]:
```

	area	pop
California	423967	38332521
Florida	170312	19552860
Illinois	149995	12882135

How to read in data

- Reading data from CSVs
 - With CSV files all you need is a single line to load in the data:

```
df = pd.read_csv('purchases.csv')
```

```
df
```

	Unnamed: 0	apples	oranges
0	June	3	0
1	Robert	2	3
2	Lily	0	7
3	David	1	2

- CSVs don't have indexes like our DataFrames, so all we need to do is just designate the **index_col** when reading:

```
df = pd.read_csv('purchases.csv', index_col=0)
```

```
df
```

	apples	oranges
June	3	0
Robert	2	3
Lily	0	7
David	1	2

How to read in data

- Reading data from JSON

```
df = pd.read_json('purchases.json')  
  
df
```

- Reading data from a SQL database

- If you're working with data from a SQL database you need to first establish a connection using an appropriate Python library, then pass a query to pandas.

	index	apples	oranges
0	June	3	0
1	Robert	2	3
2	Lily	0	7
3	David	1	2

```
df = df.set_index('index')  
  
df
```

	apples	oranges
index		
June	3	0
Robert	2	3
Lily	0	7
David	1	2

we could use `set_index()` on *any* DataFrame using *any* column at *any* time. Indexing Series and DataFrames is a very common task, and the different ways of doing it is worth remembering.

Converting back to a CSV, JSON, or SQL

- After extensive work on cleaning your data, you're now ready to save it as a file of your choice.

```
df.to_csv('new_purchases.csv')
```

```
df.to_json('new_purchases.json')
```

```
df.to_sql('new_purchases', con)
```

- When we save JSON and CSV files, all we have to input into those functions is our desired filename with the appropriate file extension.
- With SQL, we're not creating a new file but instead inserting a new table into the database using our con variable from before.

Important DataFrame Operations

- DataFrames possess hundreds of methods and other operations that are crucial to any analysis.
- As a beginner, you should know the operations that perform simple transformations of your data and those that provide fundamental statistical analysis.

Important DataFrame Operations

- Let's load in the IMDB movies dataset to begin:

```
movies_df = pd.read_csv("IMDB-Movie-Data.csv", index_col="Title")
```

- View your data

- The first thing to do when opening a new dataset is print out a few rows to keep as a visual reference.

```
movies_df.head()
```

```
movies_df.tail(2)
```

	Rank	Genre	Description	Director	Actors	Year
Title						
Search Party	999	Adventure,Comedy	A pair of friends embark on a mission to reuni...	Scot Armstrong	Adam Pally, T.J. Miller, Thomas Middleditch,Sh...	2014
Nine Lives	1000	Comedy,Family,Fantasy	A stuffy businessman finds himself trapped ins...	Barry Sonnenfeld	Kevin Spacey, Jennifer Garner, Robbie Amell,Ch...	2016

Important DataFrame Operations

- Getting info about your data
 - .info() should be one of the very first commands you run after loading your data

movies_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Index: 1000 entries, Guardians of the Galaxy to Nine Lives
Data columns (total 11 columns):
Rank                1000 non-null int64
Genre               1000 non-null object
Description         1000 non-null object
Director            1000 non-null object
Actors              1000 non-null object
Year                1000 non-null int64
Runtime (Minutes)   1000 non-null int64
Rating              1000 non-null float64
Votes              1000 non-null int64
Revenue (Millions)  872 non-null float64
Metascore           936 non-null float64
dtypes: float64(3), int64(4), object(4)
memory usage: 93.8+ KB
```

- .info() provides the essential details about your dataset, such as the number of rows and columns, the number of non-null values, what type of data is in each column, and how much memory your DataFrame is using.

Important DataFrame Operations

- Another fast and useful attribute is `.shape`, which outputs just a tuple of (rows, columns)

```
movies_df.shape
```

```
(1000, 11)
```

- Note that `.shape` has no parentheses and is a simple tuple of format (rows, columns). So we have **1000 rows** and **11 columns** in our movies DataFrame.
- You'll be going to `.shape` a lot when cleaning and transforming data. For example, you might filter some rows based on some criteria and then want to know quickly how many rows were removed.

Important DataFrame Operations

- Handling duplicates
 - This dataset does not have duplicate rows, but it is always important to verify you aren't aggregating duplicate rows.

```
temp_df = movies_df.append(movies_df)
```

```
(2000, 11)
```

```
temp_df.shape
```

```
temp_df = temp_df.drop_duplicates()
```

or

```
temp_df.drop_duplicates(inplace=True)
```

```
(1000, 11)
```

```
temp_df.shape
```

Important DataFrame Operations

- Another important argument for `drop_duplicates()` is `keep`, which has three possible options:
 - `first`: (default) Drop duplicates except for the first occurrence
 - `last`: Drop duplicates except for the last occurrence
 - `False`: Drop all duplicates

```
temp_df = movies_df.append(movies_df) # make a new copy
```

```
temp_df.drop_duplicates(inplace=True, keep=False)
```

```
temp_df.shape
```

(0, 11)

Since all rows were duplicates, `keep=False` dropped them all resulting in zero rows being left over.

Important DataFrame Operations

- Column Cleanup
 - Many times datasets will have verbose column names with symbols, upper and lowercase words, spaces, and typos.
 - To make selecting data by column name easier we can spend a little time cleaning up their names.

- Print the column names of our dataset

```
movies_df.columns
```

```
Index(['Rank', 'Genre', 'Description',  
      'Director', 'Actors', 'Year', 'Runtime  
(Minutes)', 'Rating', 'Votes', 'Revenue  
(Millions)', 'Metascore'],  
      dtype='object')
```

- Use `.rename()` method to rename certain or all columns via a dict.

```
movies_df.rename(columns={  
    'Runtime (Minutes)': 'Runtime',  
    'Revenue (Millions)':  
    'Revenue_millions'  
}, inplace=True)  
movies_df.columns
```

```
Index(['Rank', 'Genre',  
      'Description', 'Director',  
      'Actors', 'Year', 'Runtime',  
      'Rating', 'Votes',  
      'Revenue_millions',  
      'Metascore'], dtype='object') 31
```

Important DataFrame Operations

- What if we want to lowercase all names? Instead of using `.rename()` we could also set a list of names to the columns like so:

```
movies_df.columns = ['rank', 'genre',  
                     'description', 'director', 'actors',  
                     'year', 'runtime', 'rating', 'votes',  
                     'revenue_millions', 'metascore']
```

```
movies_df.columns
```

```
Index(['rank', 'genre',  
      'description', 'director',  
      'actors', 'year', 'runtime',  
      'rating', 'votes',  
      'revenue_millions',  
      'metascore'], dtype='object')
```

- But that's too much work. Instead of just renaming each column manually we can do a list comprehension:

```
movies_df.columns = [col.lower() for col  
                     in movies_df]
```

```
movies_df.columns
```

```
Index(['rank', 'genre',  
      'description', 'director',  
      'actors', 'year', 'runtime',  
      'rating', 'votes',  
      'revenue_millions',  
      'metascore'], dtype='object')
```


Work with Missing Values

- When exploring data, you'll most likely encounter missing or null values, which are essentially placeholders for non-existent values.
 - Python's **None** or NumPy's **np.nan**.
- Two options
 - Get rid of rows or columns with nulls
 - Replace nulls with non-null values, a technique known as **imputation**

Work with Missing Values

- Let's calculate the total number of nulls in each column of our dataset.

- The first step is to check which cells in our DataFrame are null:

```
movies_df.isnull()
```

	rank	genre	description	director	actors	year	runtime	rating	votes	revenue_millions	metascore
Title											
Guardians of the Galaxy	False	False	False	False	False	False	False	False	False	False	False
Prometheus	False	False	False	False	False	False	False	False	False	False	False
Split	False	False	False	False	False	False	False	False	False	False	False
Sing	False	False	False	False	False	False	False	False	False	False	False
Suicide Squad	False	False	False	False	False	False	False	False	False	False	False

`isnull()` returns a DataFrame where each cell is either True or False depending on that cell's null status.

- Then to count the number of nulls in each column we use an aggregate function for summing:

```
movies_df.isnull().sum()
```

```
rank 0 genre 0 description 0
director 0 actors 0 year 0
runtime 0 rating 0 votes 0
revenue_millions 128 metascore
64 dtype: int64
```

Work with Missing Values

- Remove null values
 - Data Scientists and Analysts regularly face the dilemma of dropping or imputing null values, and is a decision that requires intimate knowledge of your data and its context.
 - Overall, removing null data is only suggested if you have a small amount of missing data.

```
movies_df.dropna()
```

This operation will delete any **row** with at least a single null value, but it will return a new DataFrame without altering the original one. You could specify ***inplace=True*** in this method as well.

```
movies_df.dropna(axis=1)
```

Drop columns with null values

Check the reference by yourself for imputation method

Work with Missing Values

- Note that before you apply `.isnull()` and `.dropna()` methods, you need to make sure that missing values are represented properly (either `None` or `np.nan`).
- We can convert special missing values into such standard forms by specifying the `na_values` attribute when calling `pd.read_csv()`.
 - `na_values`: scalar, str, list-like, or dict, optional
 - Additional strings to recognize as NA/NaN. If dict passed, specific per-column NA values. By default the following values are interpreted as NaN: `"`, `'#N/A'`, `'#N/A N/A'`, `'#NA'`, `'-1.#IND'`, `'-1.#QNAN'`, `'-NaN'`, `'-nan'`, `'1.#IND'`, `'1.#QNAN'`, `'<NA>'`, `'N/A'`, `'NA'`, `'NULL'`, `'NaN'`, `'n/a'`, `'nan'`, `'null'`.
 - For instance: `"?"`, `"\whitespace"`, `"&"`, and `":"`.

Check Data Statistics

- Understanding your variables
 - Using `describe()` on an entire DataFrame we can get a summary of the distribution of **continuous** variables.

```
movies_df.describe()
```

	rank	year	runtime	rating	votes	revenue
count	1000.000000	1000.000000	1000.000000	1000.000000	1.000000e+03	1000.000000
mean	500.500000	2012.783000	113.172000	6.723200	1.698083e+05	82.950000
std	288.819436	3.205962	18.810908	0.945429	1.887626e+05	96.410000
min	1.000000	2006.000000	66.000000	1.900000	6.100000e+01	0.000000
25%	250.750000	2010.000000	100.000000	6.200000	3.630900e+04	17.440000
50%	500.500000	2014.000000	111.000000	6.800000	1.107990e+05	60.370000
75%	750.250000	2016.000000	123.000000	7.400000	2.399098e+05	99.170000
max	1000.000000	2016.000000	191.000000	9.000000	1.791916e+06	936.600000

- By using the correlation method `.corr()` we can generate the relationship between each continuous variable:

```
movies_df.corr()
```

	rank	year	runtime	rating	votes	revenue_millions
rank	1.000000	-0.261605	-0.221739	-0.219555	-0.283876	-0.252996
year	-0.261605	1.000000	-0.164900	-0.211219	-0.411904	-0.117562
runtime	-0.221739	-0.164900	1.000000	0.392214	0.407062	0.247834
rating	-0.219555	-0.211219	0.392214	1.000000	0.511537	0.189527
votes	-0.283876	-0.411904	0.407062	0.511537	1.000000	0.607941
revenue_millions	-0.252996	-0.117562	0.247834	0.189527	0.607941	1.000000
metascore	-0.191869	-0.079305	0.211978	0.631897	0.325684	0.133328

Check Data Statistics

- Understanding your variables

- `.describe()` can also be used on a categorical variable to get the count of rows, unique count of categories, top category, and freq of top category:

```
movies_df['genre'].describe()
```

```
count          1000
unique          207
top      Action,Adventure,Sci-Fi
freq           50
Name: genre, dtype: object
```

- `.value_counts()` can tell us the frequency of all values in a column:

```
movies_df['genre'].value_counts().head(10)
```

```
Action,Adventure,Sci-Fi    50
Drama                      48
Comedy,Drama,Romance      35
Comedy                     32
Drama,Romance              31
Action,Adventure,Fantasy   27
Comedy,Drama               27
Animation,Adventure,Comedy 27
Comedy,Romance             26
Crime,Drama,Thriller       24
Name: genre, dtype: int64
```

DataFrame Slicing, Selecting, and Extracting

- It's important to note that, although many methods are the same, DataFrames and Series have different attributes, so you'll need be sure to know which type you are working with or else you will receive attribute errors.
- By columns**

```
genre_col = movies_df['genre']
```

```
type(genre_col)
```

```
genre_col = movies_df[['genre']]
```

```
type(genre_col)
```

```
subset = movies_df[['genre', 'rating']]
```

```
subset.head()
```

```
pandas.core.series.Series
```

```
pandas.core.frame.DataFrame
```

	genre	rating
Title		
Guardians of the Galaxy	Action,Adventure,Sci-Fi	8.1
Prometheus	Adventure,Mystery,Sci-Fi	7.0
Split	Horror,Thriller	7.3
Sing	Animation,Comedy,Family	7.2
Suicide Squad	Action,Adventure,Fantasy	6.2

DataFrame Slicing, Selecting, and Extracting

- **By rows:** two options
 - `.loc` - **loc**ates by name
 - `.iloc` - **loc**ates by numerical index

Single Row

```
prom = movies_df.loc["Prometheus"]
```

1 is the numerical index of Prometheus:

```
prom = movies_df.iloc[1]
```

Multiple Rows

```
movie_subset = movies_df.loc['Prometheus':'Sing']
```

```
movie_subset = movies_df.iloc[1:4]
```

One important distinction between using `.loc` and `.iloc` to select multiple rows is that `.loc` includes the movie *Sing* in the result, but when using `.iloc` we're getting rows 1:4 but the movie at index 4 (*Suicide Squad*) is not included -- similar to Python list slicing

DataFrame Slicing, Selecting, and Extracting

- **Conditional selections**

- For example, what if we want to filter our movies DataFrame to show only films directed by Ridley Scott or films with a rating greater than or equal to 8.0?

```
condition = (movies_df['director'] == "Ridley Scott")
```

```
condition.head()
```

Title	
Guardians of the Galaxy	False
Prometheus	True
Split	False
Sing	False
Suicide Squad	False

Name: director, dtype: bool

Similar to isnull(), this returns a Series of True and False values: True for films directed by Ridley Scott and False for ones not directed by him.

```
movies_df[movies_df['director'] == "Ridley Scott"]
```

```
movies_df[movies_df['rating'] >= 8.6]
```

```
movies_df[(movies_df['director'] == 'Christopher Nolan') | (movies_df['director'] == 'Ridley Scott')]
```

DataFrame Slicing, Selecting, and Extracting

Say we want all movies that were released between 2005 and 2010, have a rating above 8.0, but made below the 25th percentile in revenue.

```
movies_df[
    ((movies_df['year'] >= 2005) & (movies_df['year'] <= 2010))
    & (movies_df['rating'] > 8.0)
    & (movies_df['revenue_millions'] < movies_df['revenue_millions'].quantile(0.25))
]
```

Apply Functions

- For example, we want to convert movies with an 8.0 or greater to a string value of "good" and the rest to "bad" and use this transformed values to create a new column.
- It is possible to iterate over a DataFrame or Series as you would with a list, but doing so — especially on large datasets — is very slow.
- Using *apply()* will be much faster than iterating manually over rows because pandas is utilizing vectorization.
 - *Vectorization: a style of computer programming where operations are applied to whole arrays instead of individual elements* — [Wikipedia](#)

Apply Functions

```
def rating_function(x):  
    if x >= 8.0:  
        return "good"  
    else:  
        return "bad"
```

```
movies_df["rating_category"] =  
movies_df["rating"].apply(rating_function)
```

The `.apply()` method passes every value in the rating column through the `rating_function` and then returns a new Series. This Series is then assigned to a new column called `rating_category`.

You can also use anonymous functions as well. This lambda function achieves the same result as `rating_function`:

```
movies_df["rating_category"] = movies_df["rating"].apply(lambda x:  
'good' if x >= 8.0 else 'bad')
```

Reading References

- Chapter 3 of [T1]