

IS 6733

Deep Learning on Cloud Platforms

Lecture 3

Deep Learning Overview

Dr. Yuanxiong Guo
Department of Information Systems and Cyber Security



Acknowledgement

- Some slides are adapted from the following course:
 - CS231n: Convolutional Neural Networks for Visual Recognition, Stanford University
- Thanks to Prof. Fei-Fei Li for sharing the slides

Deep Learning

- What is “learning”?
 - Improving performance through experience
 - Getting a computer to do well on a task without manually building in competence
- What is “deep”?
 - Successive layers of increasingly meaningful representations on the input data

Artificial Intelligence, Machine Learning, and Deep Learning

Relationship between AI, ML, and DL

ARTIFICIAL INTELLIGENCE (AI)

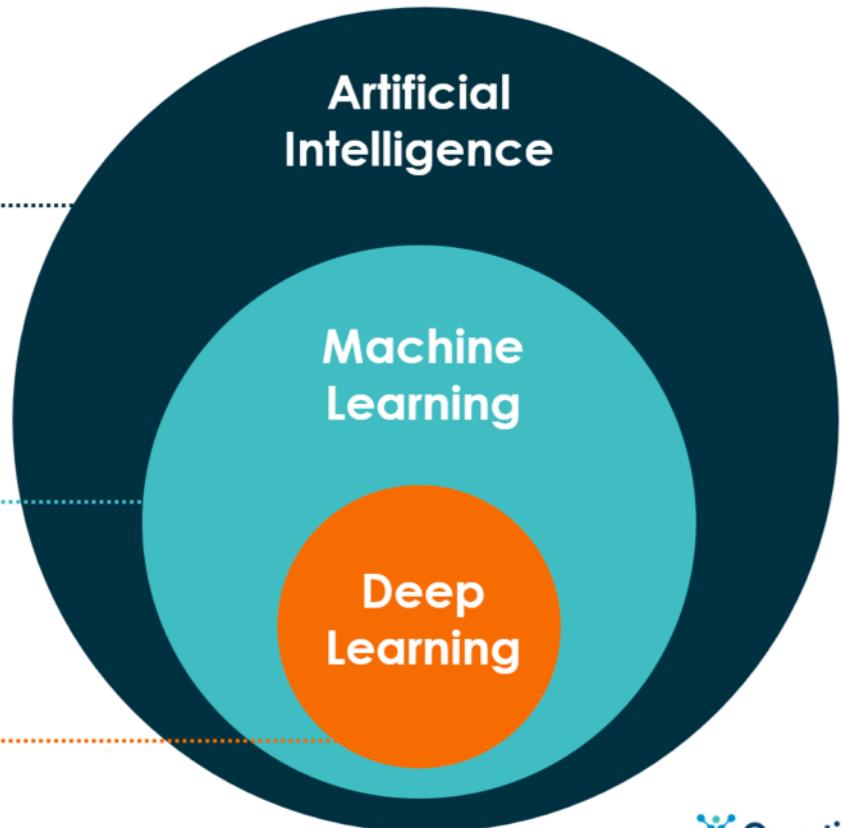
Programming systems to perform tasks which usually require human intelligence.

MACHINE LEARNING (ML)

Training algorithms to solve tasks by pattern recognition instead of specifically programming them how to solve the task.

DEEP LEARNING (DL)

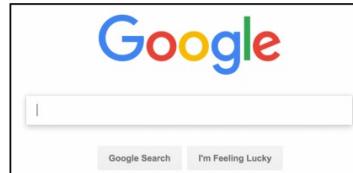
Training algorithms by using deep neural networks with multiple layers.



Today AI is ubiquitous

- Automate routine labor

- Search

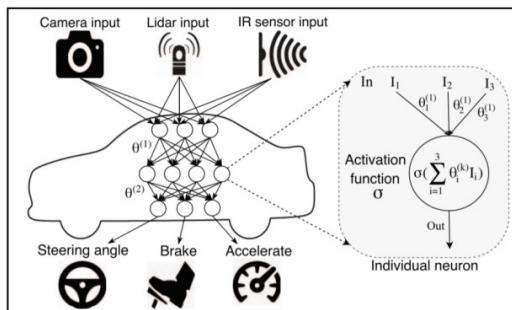


- Understand speech

- SIRI, Alexa



- Autonomous Vehicles

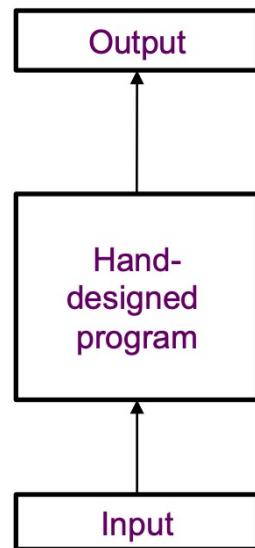


Challenge of AI

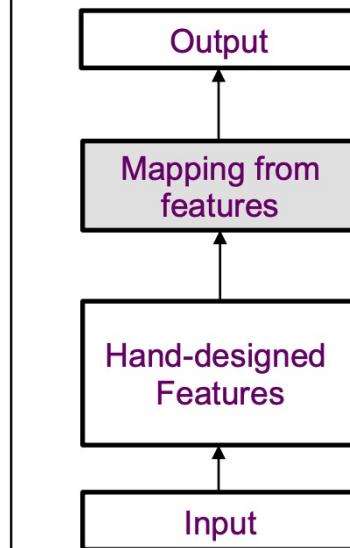
- Early successes of AI
 - Solved problems intellectually difficult for humans
 - Problems described by small set of rules
- True challenge of AI
 - Solve tasks easy for people but hard to describe formally
 - e.g., recognize spoken words, or faces in images
- Today's AI
 - It is about solving these more intuitive problems

Summary of AI Models

Rule-based System

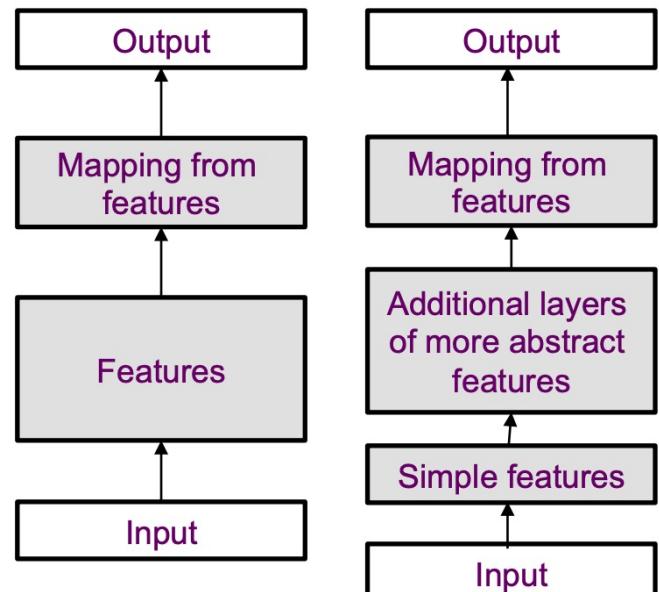


Classic Machine learning



Representation Learning

Deep Learning



Shaded boxes indicate components that can learn from data

AI Paradigm Shift

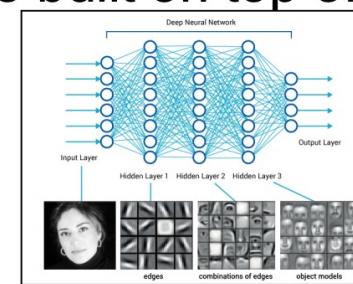
- Knowledge-based systems
 - Cannot perform simple recognition tasks
- Simple machine learning methods
 - Cannot perform complex recognition tasks
- Deep learning methods

Everyday life needs knowledge

- A person's everyday life requires immense amount of knowledge of the word
 - Much of this knowledge is intuitive and subjective
 - Difficult to articulate in a formal way
- Computers need to capture same knowledge to behave intelligently
- Key challenge of AI is how to get this informal knowledge into a computer

Solution to Intuitive Problems

1. Allow computers to learn from experience
 - By gathering knowledge from experience
 - Avoids need for human operators to specify knowledge that the computer needs
2. Understand the world as hierarchy of concepts
 - Thereby learn complicated concepts by building them out of simpler ones
 - A graph of how these concepts are built on top of each other is deep, with many layers



The Knowledge-based Approach

- Hard-coded knowledge in a formal language
 - Computer can reason about statements in these languages using inference rules
- Unwieldy process
 - Staff of human supervisors
 - People struggle formalize rules with enough complexity to describe the world

Example: Image Classification



This image by Nikita is
licensed under CC-BY 2.0

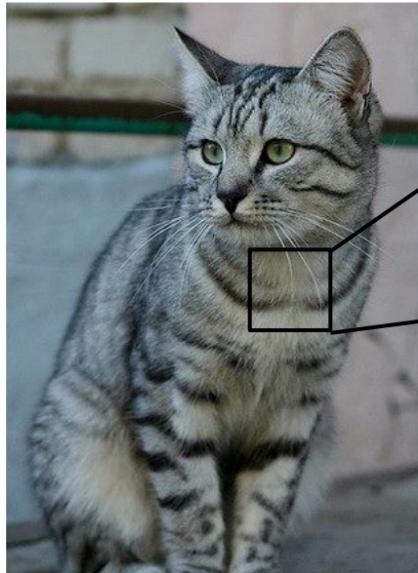
(assume given set of discrete labels)
{dog, cat, truck, plane, ...}



cat

Example: Image Classification

The Problem: Semantic Gap



This image by Nikita is
licensed under CC-BY 2.0

```
[[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]
 [ 91 98 102 106 104 79 98 103 99 105 123 136 110 105 94 85]
 [ 76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]
 [ 99 81 81 93 128 131 127 100 95 98 102 99 96 93 101 94]
 [106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]
 [114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
 [133 137 147 103 65 81 88 65 52 54 74 84 102 93 85 82]
 [128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]
 [125 133 144 137 119 121 117 94 65 79 80 65 54 64 72 98]
 [127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
 [115 114 109 123 150 148 131 118 113 109 100 92 74 65 72 78]
 [ 89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 80]
 [ 63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]
 [ 62 65 82 89 78 71 88 101 124 126 119 101 107 114 131 119]
 [ 63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]
 [ 87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
 [118 97 82 86 117 123 116 66 41 51 95 93 89 95 102 107]
 [164 146 112 80 82 120 124 104 76 48 45 66 89 181 182 109]
 [157 170 157 120 93 86 114 132 112 97 69 55 78 82 99 94]
 [130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]
 [128 112 96 117 150 144 128 115 104 107 102 93 87 81 72 79]
 [123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]
 [122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]
 [122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]
```

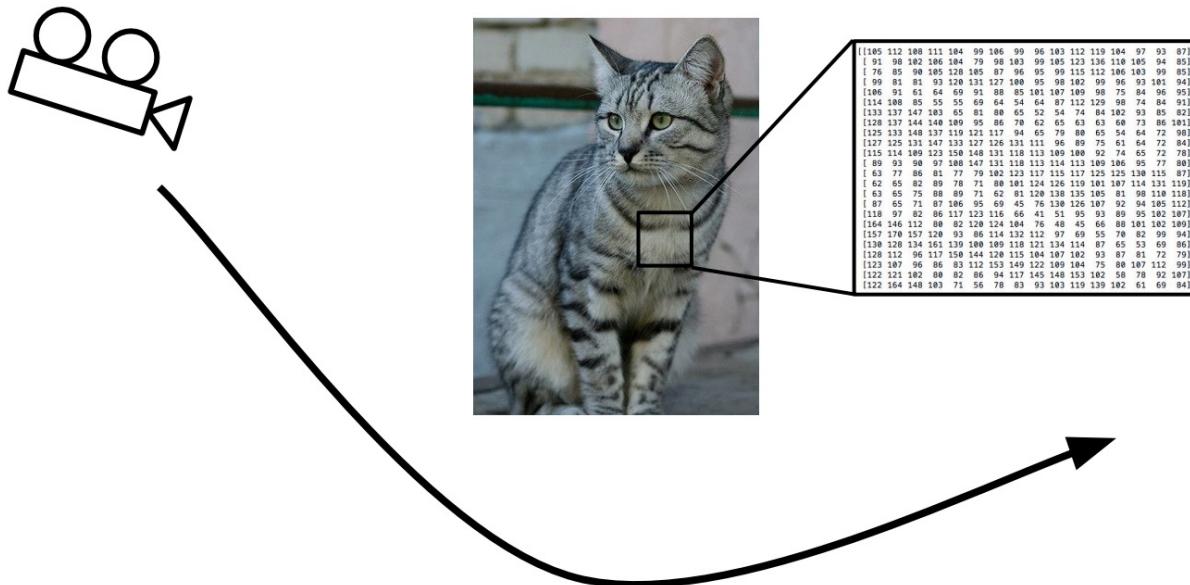
What the computer sees

An image is just a big grid of numbers between [0, 255]:

e.g. 800 x 600 x 3
(3 channels RGB)

Example: Image Classification

Challenges: Viewpoint variation



Example: Image Classification

Challenges: Background Clutter



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

Example: Image Classification

Challenges: Illumination



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

Example: Image Classification

Challenges: Deformation



This image by [Umberto Salvagnin](#)
is licensed under [CC-BY 2.0](#)



This image by [Umberto Salvagnin](#)
is licensed under [CC-BY 2.0](#)



This image by [sare bear](#) is
licensed under [CC-BY 2.0](#)



This image by [Tom Thai](#) is
licensed under [CC-BY 2.0](#)

Example: Image Classification

Challenges: Occlusion



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



[This image by **jonsson** is licensed under CC-BY 2.0](#)

Example: Image Classification

Challenges: Intraclass variation



This image is CC0 1.0 public domain

Example: Image Classification

An image classifier

```
def classify_image(image):
    # Some magic here?
    return class_label
```

Unlike e.g. sorting a list of numbers,

no obvious way to hard-code the algorithm for
recognizing a cat, or other classes.

The ML Approach

- Difficulties of hand-coded approach suggests:
 - AI systems need ability to acquire knowledge by extracting patterns from raw data
- This approach is known as machine learning
 - It allowed tackling problems requiring knowledge of real world and make decisions that seem subjective
 - Simple ML algorithms
 - Logistic Regression
 - Naive Bayes

The ML Approach

Machine Learning: Data-Driven Approach

1. Collect a dataset of images and labels
2. Use Machine Learning to train a classifier
3. Evaluate the classifier on new images

```
def train(images, labels):
    # Machine learning!
    return model
```

```
def predict(model, test_images):
    # Use model to predict labels
    return test_labels
```

Example training set

airplane



automobile



bird



cat



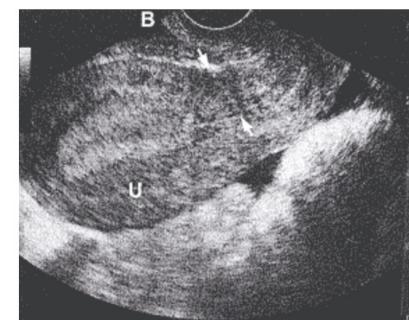
deer



Simple ML depends on representation

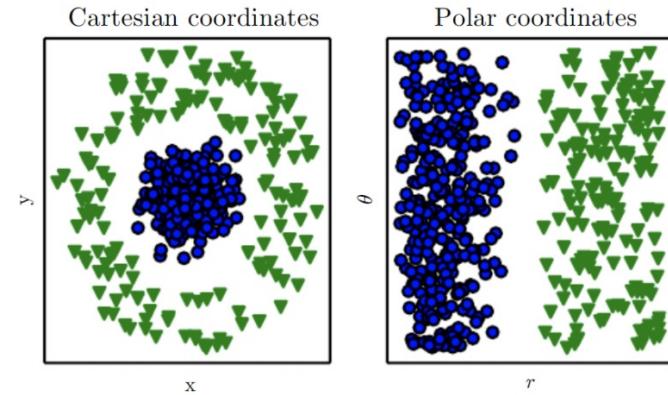
- Simple ML algorithms depend heavily on *representation* of given data
- Ex: Logistic regression to recommend delivery does not examine patient directly
 - Instead doctor provides information such presence of uterine scar (called a feature)
 - If logistic regression was provided with MRI scan, it could not make useful predictions

Individual pixels have negligible correlation with complications in delivery



Dependence on Representation

- This dependence appears in computer science and daily life
 - In CS: search is exponentially faster if data is structured and indexed intelligently
 - People: arithmetic on Arabic numerals 1,2,...9,10, easier than using Roman numerals I, II,...IX, X
 - Choice of representation has enormous effect on ML algorithm performance
 - Ex: straight line separation
 - Impossible in Cartesian
 - Simple straight line in polar



Designing right set of features

- Difficult to know what set of features are good for detecting a car in photographs
- Presence of wheel as a feature
 - Has a simple geometric shape but difficult to describe in terms of pixel values
 - Shadows, glare from metal parts, fender or object in front obscures

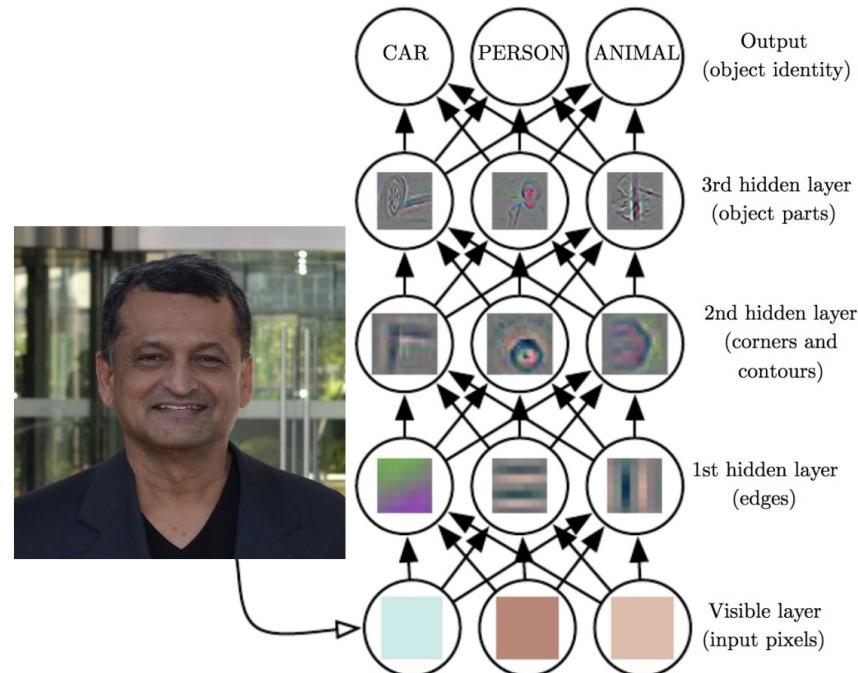


Representation Learning

- Solution: use ML to **not only learn mapping from representation to output but representation itself**
 - Better results than hard-coded representations
- Allows AI systems to rapidly adapt to new tasks
 - Designing features can take great human effort
 - Can take decades for a community of researchers

Feature Learning for Classification

- Function to map pixels to object identity is complicated
- Series of hidden layers extract increasingly abstract features
- Final decision made by a simple classifier



Characteristics of Deep Learning

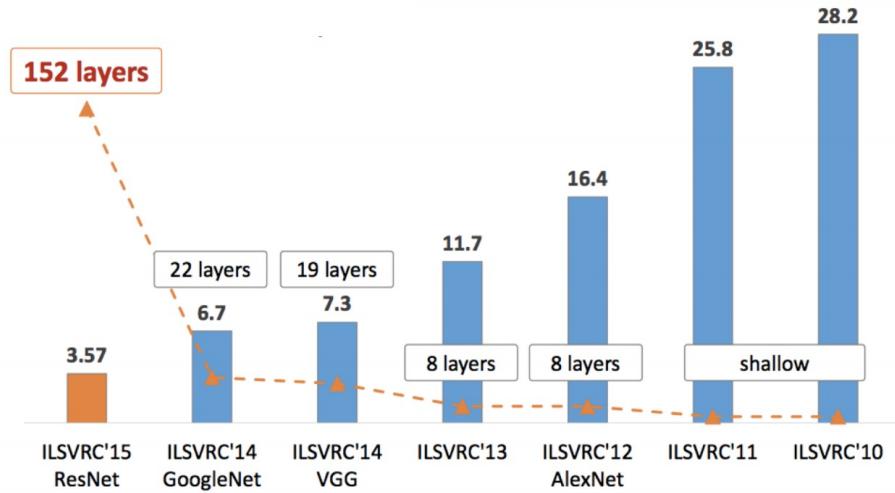
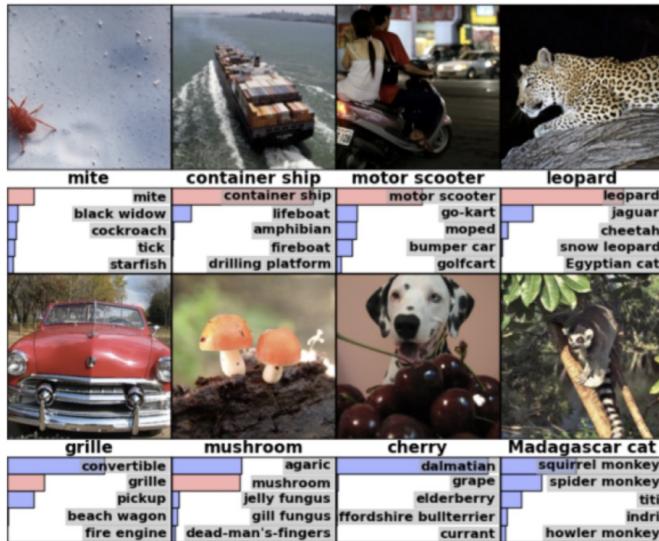
- A type of Machine Learning that improves with experience and data
- Only viable approach to building AI systems in real-world environments
- Obtains its power by a nested hierarchy of concepts
 - each concept defined by relationship to simpler concepts
 - More abstract representations computed in terms of less abstract ones

Success of Deep Learning

- Vision
- Speech and language
- Game
- Robotics
- Self-driving cars

Vision

- ImageNet Challenge



Vision

- Object detection, instance segmentation



K. He, G. Gkioxari, P. Dollar, and R. Girshick, [Mask R-CNN](#),
ICCV 2017 (Best Paper Award)

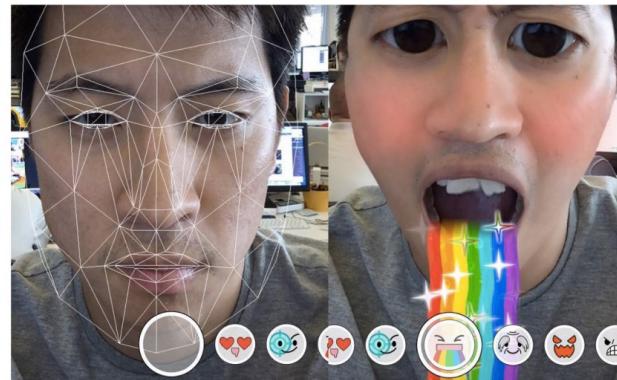
Vision



[Facebook accessibility tools for the visually impaired](#)



[AI beats human pathologists at detecting cancer](#)



[Technology behind Snapchat lenses](#)

Vision



[Beijing bets on facial recognition in a big drive for total surveillance](#) – Washington Post, 1/8/2018

Speech and natural language



Skype Translator

Break down the language barrier with your friends, family and colleagues.

Our online translator can help you communicate in 7 languages for voice calls, and in more than 50 languages while instant messaging.

Skype Translator uses machine learning. So the more you use it, the better it gets. Thanks for being patient as the technology graduates from Preview mode.

<https://www.skype.com/en/features/skype-translator/>



Google Translate App

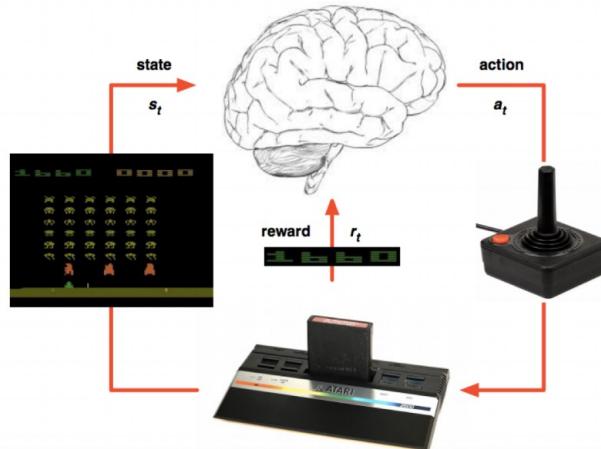
- Translate between 103 languages by typing
- Offline: Translate 52 languages when you have no Internet
- Instant camera translation: **Use your camera to translate text** instantly in 30 languages
- Camera Mode: Take pictures of text for higher-quality translations in 37 languages
- Conversation Mode: Two-way **instant speech translation** in 32 languages

<https://play.google.com/store/apps/details?id=com.google.android.apps.translate&hl=en>

See also: [The Great AI Awakening](#) (New York Times Magazine)

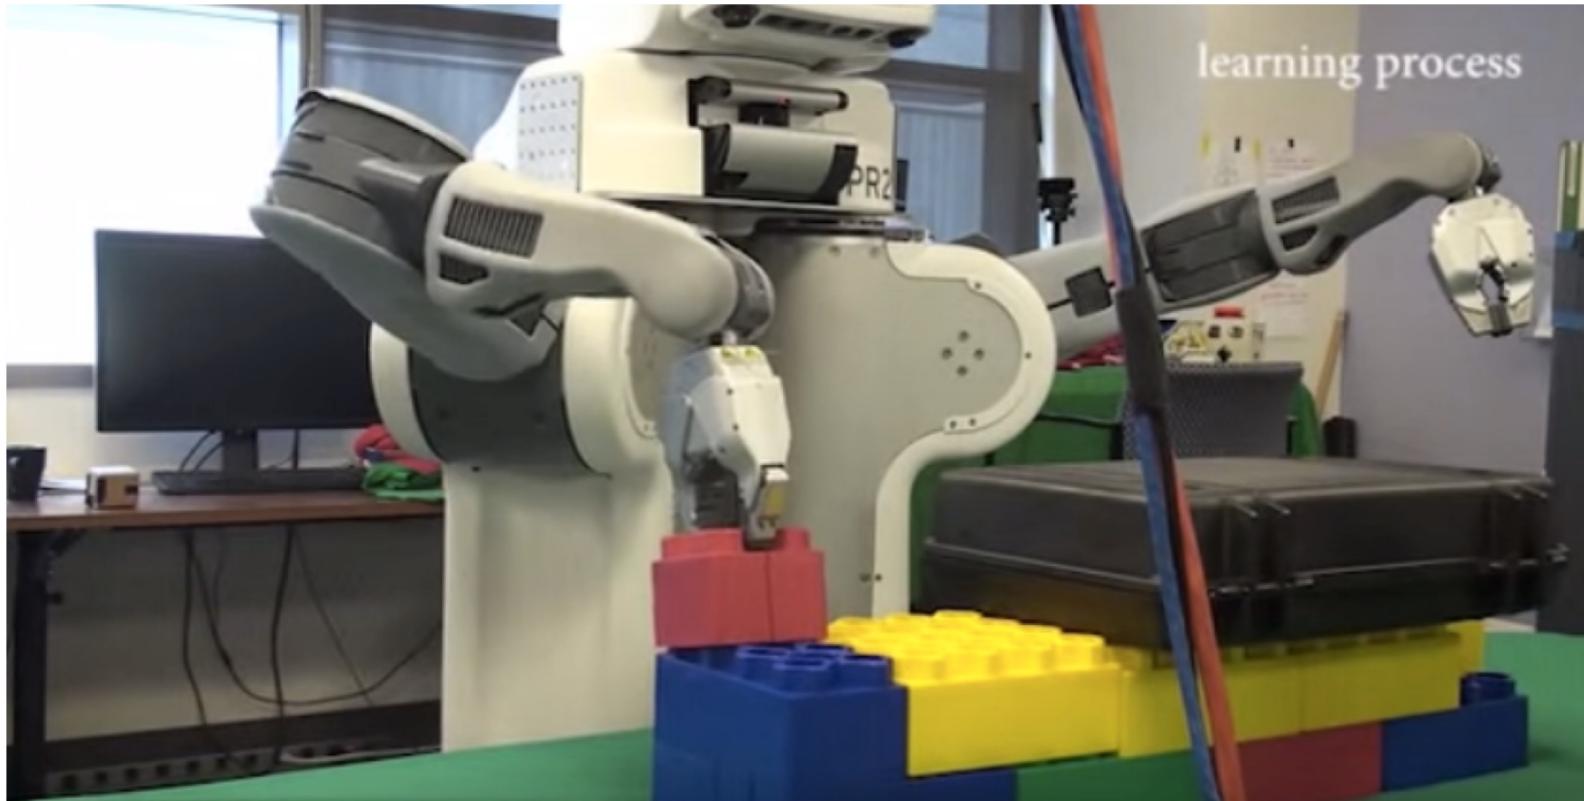
Games

- 2013: DeepMind uses deep reinforcement learning to beat humans at some Atari games
- 2016: DeepMind's AlphaGo system beats Go grandmaster Lee Sedol 4-1
- 2017: AlphaZero learns to play Go and chess from scratch



Robots

- End-to-end training of deep visuomotor policies



<https://www.youtube.com/watch?v=Q4bMcUk6pcw>

Self-driving cars

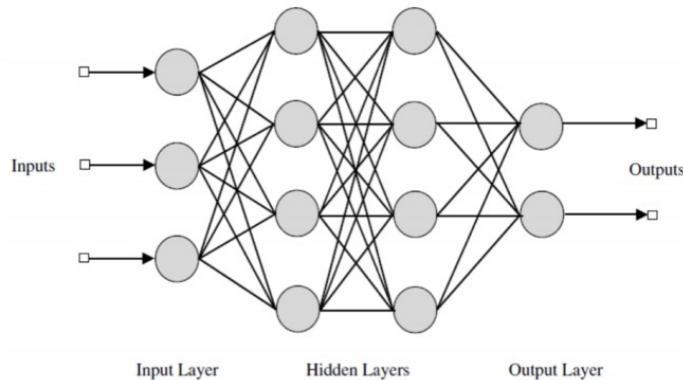


[Image source](#)

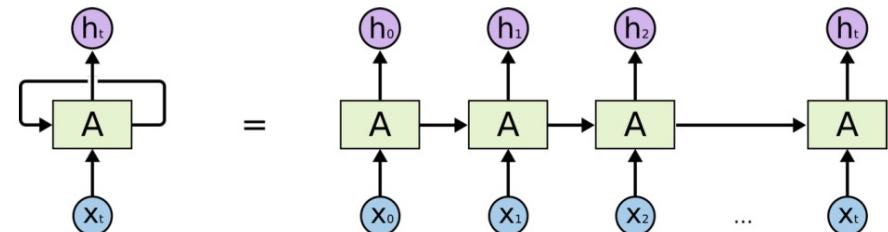
- Deep learning crucial for the global success of automotive autonomy – [Automotive World, 6/26/2018](#)

Basic Deep Learning Models

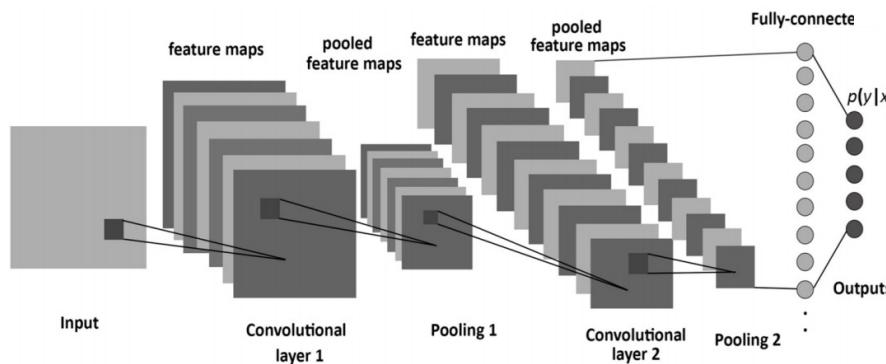
Multilayer neural networks, backpropagation



Recurrent networks and applications



Convolutional networks and applications



Before DL: A Brief History of ML

A Few Facts about DL

- DL has reached a level of public attention and industry investment never seen before in the history of AI

But,

- DL isn't the first successful form of ML
- Most of the ML algorithms used in the industry today aren't DL algorithms.
- DL isn't always the right tool for the job
 - Not enough data for DL to be applicable
 - Problems better solved by a different ML algorithm

Probabilistic Modeling

- The application of the principles of statistics to data analysis
- One of the earliest forms of ML and still widely used to this day
 - Naïve Bayes
 - Classifiers based on applying Bayes' theorem while assuming the features in the input data are all independent (a strong, or “naïve” assumption)
 - Logistic Regression
 - Often the first thing a data scientist will try on a dataset to get a feel for the classification task at hand

Early Neural Networks

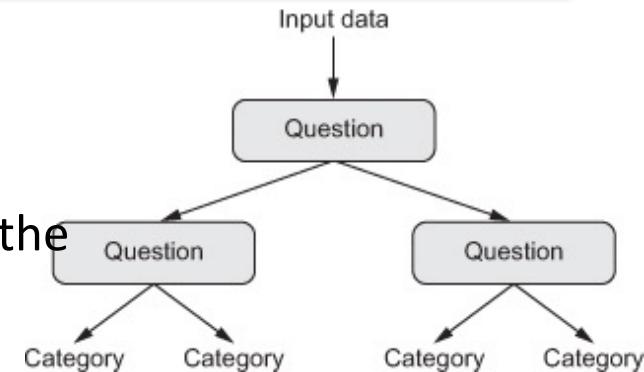
- The core ideas of neural networks were investigated in toy forms as early as 1950s
- The major difficulty at that time was an efficient way to train large neural networks
- Changed in mid-1980s due to the invent of *backpropagation* algorithm
- The first successful application of neural nets came in 1980 from Bell Labs
 - Yann LeCun applied CNN and backpropagation to classify handwritten digits
 - The resulting network (LeNet) was used by USPS in the 1990s to automate the reading of ZIP code on mail envelopes

Kernel Methods

- As neural networks started to gain some attention in the 1990s, *Kernel methods* rose to fame and quickly sent neural nets back to oblivion.
- *Kernel methods*: a group of classification algorithms where
 1. Data is mapped to a new high-dimensional representation where the decision boundary can be expressed as a hyperplane (using kernel function to efficiently compute the distance between pairs of points in the new representation)
 2. A good decision boundary is computed by trying to maximize the distance between the hyperplane and the closest data points from each class
- SVMs exhibited state-of-the-art performance on simple classification problems with rigorous theory support
 - Did not provide good results for perceptual problems such as image classification (shallow method)
 - Requires feature engineering before being applied

Decision Trees, Random Forests, and Gradient Boost Machines

- *Decision trees* are flowchart-like structures
 - Easy to visualize and interpret
 - Began to receive significant research interest in the 2000s, and by 2010 often preferred to kernel methods
- *Random forest* involves building a large number of specialized decision trees and then ensembling their outputs
 - Applicable to a wide range of problems
 - Almost always the second-best algorithm for any shallow ML task
- *Gradient boosting machines*: similar to a random forest but uses gradient boosting
 - A way to improve any ML model by iteratively training new models that specialize in addressing the weak points of the previous models
 - One of the best algorithm for dealing with nonperpetual data today



Back to Neural Networks

- Around 2010, neural networks were almost completely shunned by the scientific community at large
- The watershed moment came in 2012: Geoffrey Hinton at UToronto achieved 83.6% vs. 74.3% in ImageNet challenge using deep convolutional neural networks.
- Since 2012, deep convolutional neural networks have become the go-to algorithm for all computer vision tasks, or more generally, all perceptual tasks.

The Modern Machine Learning Landscape

- Two dominant approaches used by practitioners:
gradient boosting machines and *deep learning*
 - *Gradient boosting* is used for problems where structured data is available (shallow-learning problems)
 - XGBoost library (Python and R)
 - *Deep learning* is used for perceptual problems such as image classification
 - Keras library (Python and R)

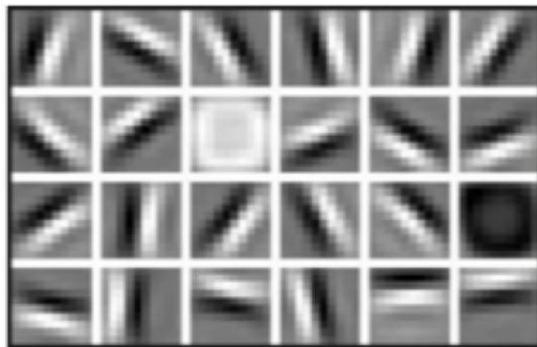
Why Deep Learning? Why Now?

Why Deep Learning?

Hand engineered features are time consuming, brittle and not scalable in practice

Can we learn the **underlying features** directly from data?

Low Level Features



Lines & Edges

Mid Level Features



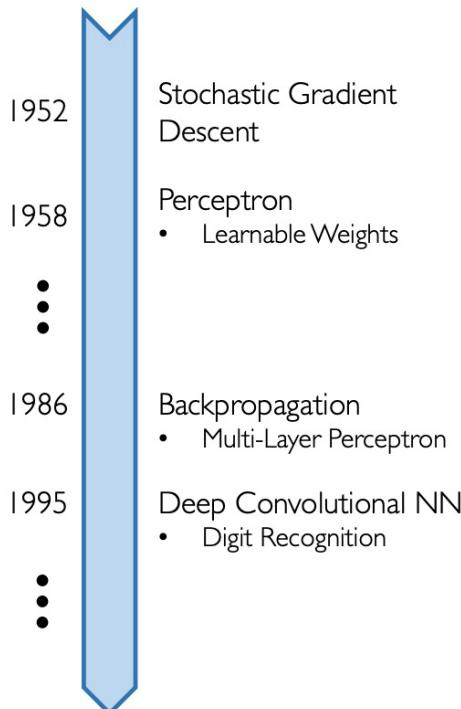
Eyes & Nose & Ears

High Level Features



Facial Structure

Why Now?



Neural Networks date back decades, so why the resurgence?

1. Big Data

- Larger Datasets
- Easier Collection & Storage



WIKIPEDIA
The Free Encyclopedia



2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable



3. Software

- Improved Techniques
- New Models
- Toolboxes



Deep Learning Hardware

Inside a computer



Spot the CPU!

(central processing unit)



This image is licensed under [CC-BY 2.0](#)



Spot the GPUs!

(graphics processing unit)



[This image](#) is in the public domain



NVIDIA

vs

AMD

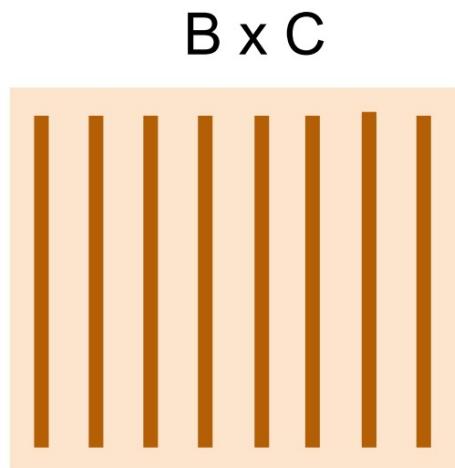
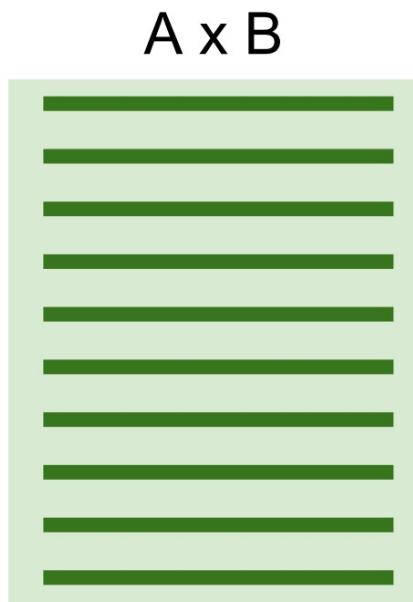
CPU vs GPU

	Cores	Clock Speed	Memory	Price	Speed
CPU (Intel Core i7-7700k)	4 (8 threads with hyperthreading)	4.2 GHz	System RAM	\$385	~540 GFLOPs FP32
GPU (NVIDIA RTX 2080 Ti)	3584	1.6 GHz	11 GB GDDR6	\$1199	~13.4 TFLOPs FP32

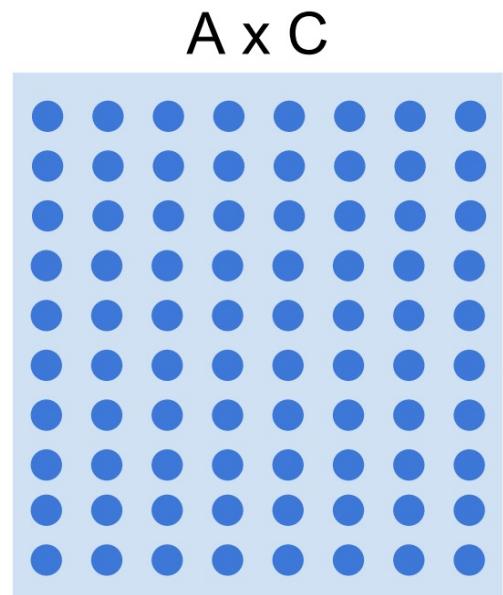
CPU: Fewer cores, but each core is much faster and much more capable; great at sequential tasks

GPU: More cores, but each core is much slower and “dumber”; great for parallel tasks

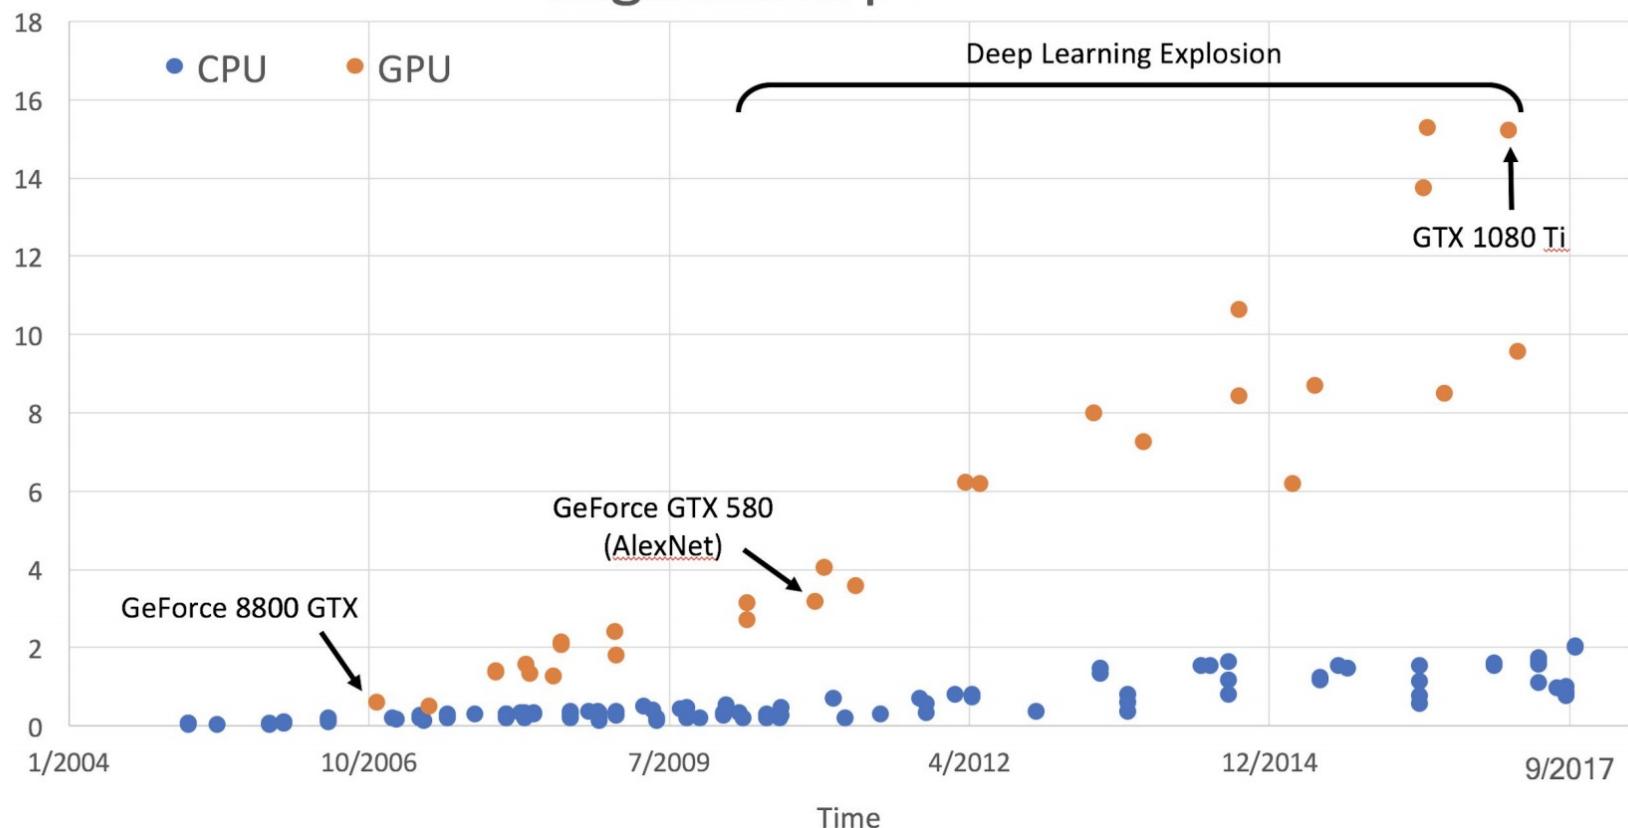
Example: Matrix Multiplication



=

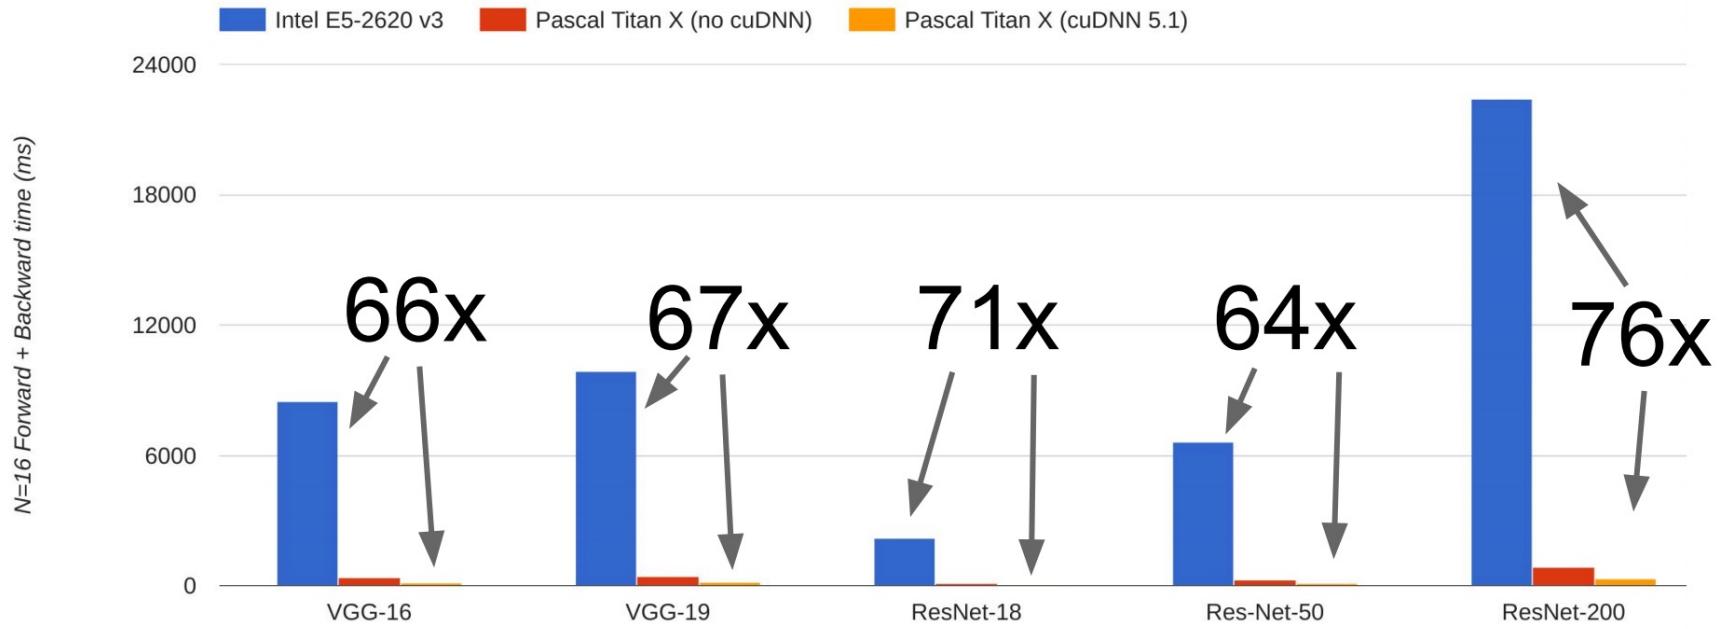


GigaFLOPs per Dollar



CPU vs GPU in practice

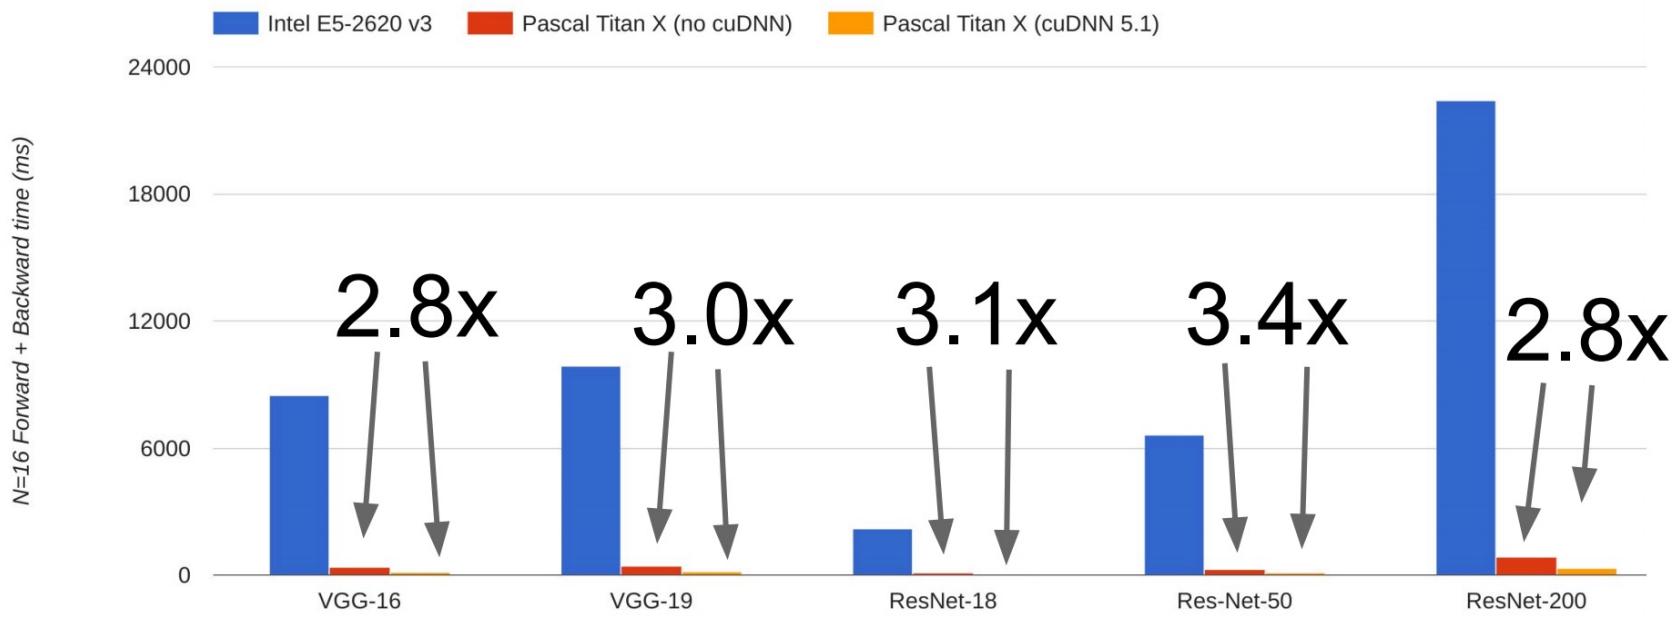
(CPU performance not well-optimized, a little unfair)



Data from <https://github.com/cjohnson/cnn-benchmarks>

CPU vs GPU in practice

cuDNN much faster than
“unoptimized” CUDA



Data from <https://github.com/jcjohnson/cnn-benchmarks>

NOTE: TITAN V isn't technically a "TPU" since that's a Google term, but both have hardware specialized for deep learning

CPU vs GPU

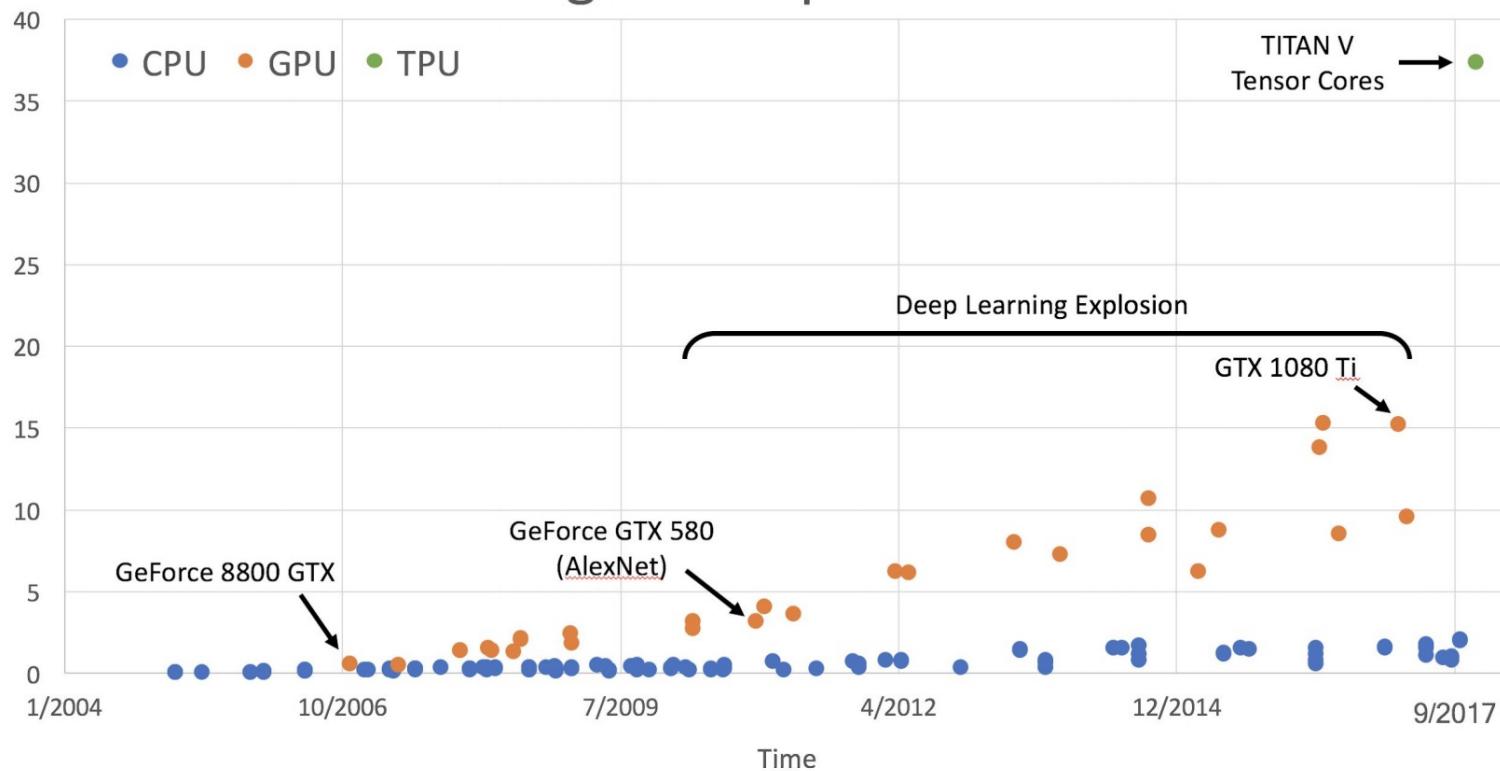
	Cores	Clock Speed	Memory	Price	Speed
CPU (Intel Core i7-7700k)	4 (8 threads with hyperthreading)	4.2 GHz	System RAM	\$385	~540 GFLOPs FP32
GPU (NVIDIA RTX 2080 Ti)	3584	1.6 GHz	11 GB GDDR6	\$1199	~13.4 TFLOPs FP32
TPU NVIDIA TITAN V	5120 CUDA, 640 Tensor	1.5 GHz	12GB HBM2	\$2999	~14 TFLOPs FP32 ~112 TFLOP FP16
TPU Google Cloud TPU	?	?	64 GB HBM	\$4.50 per hour	~180 TFLOP

CPU: Fewer cores, but each core is much faster and much more capable; great at sequential tasks

GPU: More cores, but each core is much slower and "dumber"; great for parallel tasks

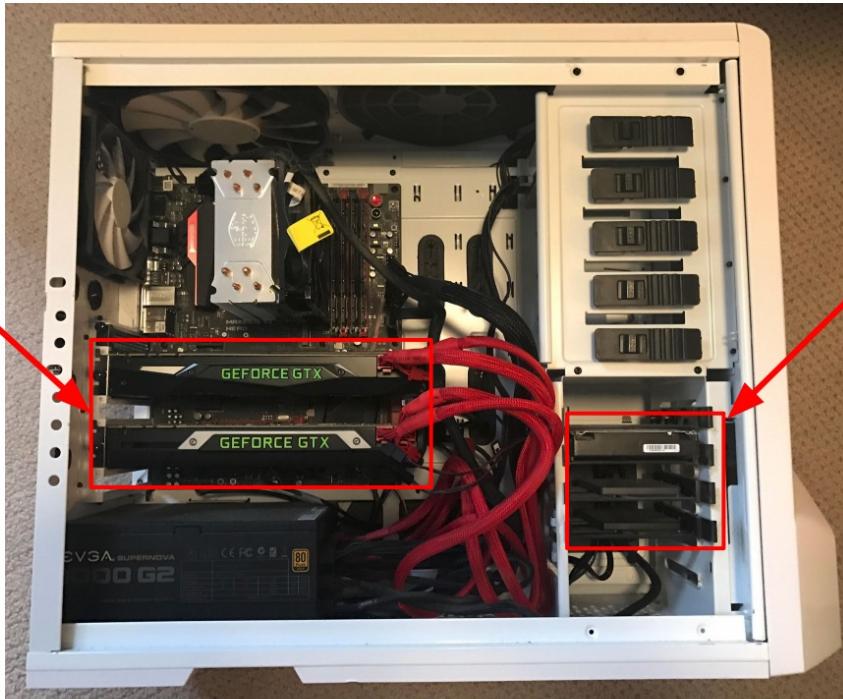
TPU: Specialized hardware for deep learning

GigaFLOPs per Dollar



CPU / GPU Communication

Model
is here



Data is here

If you aren't careful, training can bottleneck on reading data and transferring to GPU!

Solutions:

- Read all data into RAM
- Use SSD instead of HDD
- Use multiple CPU threads to prefetch data

Deep Learning Software

A zoo of frameworks!

Caffe
(UC Berkeley)



Torch
(NYU / Facebook)



Theano
(U Montreal)



Caffe2
(Facebook)

PyTorch
(Facebook)

TensorFlow
(Google)

PaddlePaddle
(Baidu)

MXNet
(Amazon)

Developed by U Washington, CMU, MIT, Hong Kong U, etc but main framework of choice at AWS

Chainer

CNTK
(Microsoft)

JAX
(Google)

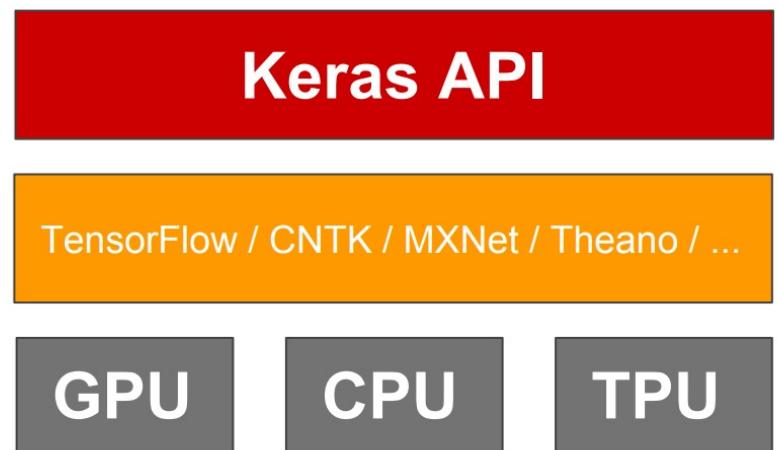
And others...

The point of deep learning framework

- Quick to develop and test new ideas
- Automatically compute gradients
- Run it all efficiently on GPU (wrap cuDNN, cuBLAS, etc)

Keras is multi-backend, multi-platform

- Develop in Python, R
 - On Unix, Windows, OSX
- Run the same code with...
 - TensorFlow
 - CNTK
 - Theano
 - MXNet
 - PlaidML
 - - ??
- CPU, NVIDIA GPU, AMD GPU, TPU...



Our Focus – Keras API of TensorFlow

- Keras is the official high-level API of TensorFlow

- tensorflow.keras (tf.keras) module
- Part of core TensorFlow since v1.4
- Full Keras API
- Better optimized for TF
- Better integration with TF-specific features
 - Estimator API
 - Eager execution
 - etc.



```
# import tf.keras as part of your TensorFlow program setup:  
import tensorflow as tf  
from tensorflow import keras
```

Reading Materials

- Chapter 1 of [T2]